

CIS 2033, Spring 2017

Instructor: David Dobor

April 23, 2017

We consider a very practical application of the weak law of large numbers and of the central limit theorem. The application has to do with polling.

The polling problem

There's a certain referendum that's going to take place. We're close enough to the day of the referendum so that voters have made up their minds, and there is a fraction p of the population that represents the voters that are going to vote "yes". But the referendum has not yet taken place, and you want to find out, to predict or estimate, what p actually is.

What you do is that you go ahead, and you select, at random, a number of people out of the population. For each person, you record their answer, whether they intend to vote "yes", or whether they intend to vote "no", and you let

$$X_i = \begin{cases} 1 & \text{if the } i\text{-th selected person says "yes"} \\ 0 & \text{if the } i\text{-th selected person says "no"} \end{cases}$$

When we say that the people are randomly selected, what we mean is that we choose them uniformly from the population. And since there's a fraction p that will vote yes, this means that the random variable X_i will be 1 with probability p , and therefore the expected value of X_i is equal to p .

In addition, we assume that we select those people independently. Now note that if we select people independently, there's always a chance that the first person polled will be the same as the second person polled, something that we do not really want to happen. However, if we assume that the population is very large, or even idealize the situation by assuming that the population is infinite, then this is never going to happen, and this will not be a concern.

So how do we proceed? We look at the results that we get from the people that we polled. We count how many said "yes", divide that count by n , and this gives us the fraction of "yes"es *in the sample*:

$$M_n = \frac{X_1 + \dots + X_n}{n}.$$

p is the fraction of the population that will vote "yes" in a referendum.

Since X_i is a Bernoulli random variable with success probability p , its mean is

$$\mathbb{E}[X_i] = p,$$

and its variance is

$$\sigma_{X_i}^2 = p(1-p).$$

M_n is the fraction of the people in the sample who voted "yes".

And M_n is a pretty reasonable estimate for the unknown fraction p , the fraction of "yes"es in the overall population.

Now perhaps your boss has asked you to find out the exact value of p . What should your response be? Well, there is *no way* to calculate p exactly on the basis of a finite and random poll.

Therefore, there is going to be some error in our estimation of p .

Then, perhaps your boss comes back and says: "OK, then try to give me an estimate of p which is very accurate. I would like you to come up with an estimate which is correct within one percentage point. Can you do this for me?"

Your answer might be: "OK, let me try polling 10,000 people, and see if I can guarantee for you such a small error."

But after you think about the situation a little more, you realize that there is no way of guaranteeing such a small error with certainty. What if you are really unlucky, and the people that you poll happen to be not representative of the true voting population?

So you come back to your boss and say: "I cannot guarantee with certainty that the error is going to be small, but perhaps I can guarantee for you that the error that I get is small with high probability." Or, alternatively, "I'm going to guarantee for you that the probability that we get an error that's bigger than 0.01 is very small:

$$\mathbf{P}(|M_{10,000} - p| \geq 0.01) \leq \text{"something small"} \quad (1)$$

So how small is it going to be? Let's try to derive a bound on this probability of an error larger than one percentage point.

Using the weak law of large numbers to find the "something small" in equation (1)

First, recall what exactly were the calculations that we carried out when we derived the weak law of large numbers. We saw that the weak law of large numbers follows directly from Chebyshev's inequality that says that for any random variable X with finite mean μ and variance σ^2 , it holds that

$$\mathbf{P}(|X - \mu| \geq \epsilon) \leq \frac{\sigma^2}{\epsilon^2}, \quad \text{for any } \epsilon > 0. \quad (2)$$

Now, let our M_n play the role of X in Chebyshev's inequality (2). To rewrite (2) for our particular purpose, we need the variance of M_n and we need its mean. We computed these quantities earlier and found that $\mathbf{E}[M_n] = p$, and $\sigma_{M_n}^2 = \frac{p(1-p)}{n}$.

Your boss would like a "small error", e.g., within one percentage point: $|M_n - p| < 0.01$.

There is no way of guaranteeing *with certainty* that the error $|M_n - p|$ will be within one percentage point (i.e., will be < 0.01).

However, you can give a guarantee in the form of equation (1), which says that "the chances that the fraction of "yes"es in your sample, M_n , deviating from the true fraction of the population of voters that will vote "yes", p , can be made "small".

The statement in (2) is the familiar Chebyshev inequality.

So, using the calculations that we carried out when we derived the weak law of large numbers, we know that

$$\mathbf{P}(|M_{10,000} - p| \geq 0.01) \leq \frac{\sigma^2}{10000 \cdot (0.01)^2} = \frac{p(1-p)}{10^4 \cdot 10^{-4}} = p(1-p)$$

OK, but now, the numerator of the above expression depends on p , and we do not know what p is. (Of course if we knew p , there would be no need to conduct any polling in the first place.)

However, if you take the expression $p(1-p)$ and plot it as a function of p , what you obtain is a plot you see in Figure 1. The maximum happens when p is equal to $1/2$, in which case we get a value of $1/4$.

That is, the variance of the Bernoulli is, at most, $1/4$. Therefore, rewriting the previous equation, we obtain the following bound:

$$\mathbf{P}(|M_{10,000} - p| \geq 0.01) \leq p(1-p) \leq \frac{1}{4}$$

So you tell your boss: "If I sample 10,000 people, then the probability of an error more than the one percentage point is going to be less than 25%."

At which point, your boss might reply and say: "Well, a probability of a large error of 25% is just too big. This is unacceptable. I would like you to have a probability of error that's less than 5%."

So suppose now that we want to reduce the error, and make it only 5% or less. How are you going to proceed? You take another look at Chebyshev's inequality and rewrite it for any n , not just for $n = 10000$:

$$\mathbf{P}(|M_n - p| \geq 0.01) \leq \frac{\sigma^2}{n \cdot (0.01)^2} = \frac{p(1-p)}{n \cdot 10^{-4}} \leq \frac{1}{4 \cdot n \cdot 10^{-4}}.$$

So if we want to make

$$\frac{1}{4 \cdot n \cdot 10^{-4}} \leq 0.05 = \frac{5}{10^2}, \quad (3)$$

we solve (3) for n to obtain that

$$n \geq \frac{10^6}{20} = 50,000. \quad (4)$$

So at this point, you go back to your boss and tell them that one way of guaranteeing that the probability of a large error is less than or equal to 5% is to take n equal to 50,000, and that this number of 50,000 will suffice to achieve the desired specs:

Since M_n is the sum on n independent Bernoulli(p) random variables divided by n , the mean of M_n is

$$\mathbf{E}[M_n] = \frac{np}{n} = p,$$

and its variance is

$$\sigma_{M_n}^2 = \frac{np(1-p)}{n^2} = \frac{p(1-p)}{n}.$$

In particular, since we take $n = 10,000$, we have that

$$\sigma_{M_{10,000}}^2 = \frac{p(1-p)}{10^4}.$$



Figure 1: A plot of $\sigma = p(1-p)$. The maximum occurs when $p = 1/2$, at which point $\sigma = 1/4$.

$$\mathbf{P}(|M_{50,000} - p| \geq 0.01) \leq 0.05$$

Notice that the desired specs have two parameters. One is the **accuracy, 0.01**, that you want, and the other reflects the **confidence, 0.05**, with which the accuracy is going to be achieved.

Now 50,000 is a pretty large number. If you notice the results of polls, when they are presented in newspapers or on websites, they usually tell you that there's an accuracy of plus or minus *three percentage points*, not one percentage point. That helps things a little: it means that you can do with a somewhat smaller sample size.

And then there's another effect. Our calculations were based on the Chebyshev inequality. But the Chebyshev inequality is not that accurate. It turns out that if we use more accurate estimates of the probability $\mathbf{P}(|M_n - p| \geq \epsilon)$, we will find that actually much smaller values of n will be enough for our purposes.

So you can say that you are at least **95% confident** that the sample mean is within **1 percentage point** of the true mean, p , when $n = 50,000$ people are sampled.

Applying the Central Limit Theorem to Polling

We will now use the central limit theorem to approximate the quantity $\mathbf{P}(|M_n - p| \geq 0.01)$.

Recall that the central limit theorem involves the standardized version of the random variable S_n , where S_n stands for the sum of the X_i 's:

$$Z_n = \frac{S_n - n\mu}{\sqrt{n}\sigma}, \quad (5)$$

and we know that the random variable Z_n is approximately normal for sufficiently large n .

WHAT WE want to do now is to take the event

$$|M_n - p| \geq 0.01 \quad (6)$$

and rewrite it in an equivalent way, but in a way that involves the random variable Z_n .

FIRST, we note that in (5) we have a μ and a σ , so we should keep in mind what these are. For a Bernoulli random variable, the mean is p and the standard deviation σ is $\sqrt{p(1-p)}$.

$$\begin{aligned} \mu &\text{ is } p \\ \sigma &\text{ is } \sqrt{p(1-p)} \end{aligned}$$

NEXT, let's look at the event in (6). M_n is S_n/n and so we can say that

$$|M_n - p| \geq 0.01 \text{ is the same as } \left| \frac{S_n}{n} - \frac{np}{n} \right| \geq 0.01,$$

which is the same as

$$\left| \frac{S_n - np}{n} \right| \geq 0.01. \quad (7)$$

So (7) and (6) are the same event, and (7) starts to look like the expression in (5): p is the same as μ , but there is a little bit of a difference in the denominator terms. So let's see what we can do. Let's take the event in (7) but multiply both sides of the inequality by \sqrt{n} . This causes the denominator term to become just \sqrt{n} and we get a \sqrt{n} term in the numerator on the other side:

$$\left| \frac{S_n - np}{n} \right| \geq 0.01 \text{ is the same as } \left| \frac{S_n - np}{\sqrt{n}} \right| \geq 0.01\sqrt{n}$$

Let $S_n = X_1 + \dots + X_n$. Then, by the central limit theorem, we know that

$$Z_n = \frac{S_n - n\mu}{\sqrt{n}\sigma}$$

is approximately standard normal for large enough n .

Finally, we multiply the denominators on both sides by σ and we obtain this equivalent representation:

$$\left| \frac{S_n - np}{\sqrt{n}} \right| \geq 0.01\sqrt{n} \text{ is the same as } \left| \frac{S_n - np}{\sqrt{n}\sigma} \right| \geq \frac{0.01\sqrt{n}}{\sigma}. \quad (8)$$

But now we notice that in (8) we do have the random variable Z_n that we wanted, and so we managed to express the event $|M_n - p| \geq 0.01$ in terms of the random variable Z_n . In short, what we have is that

$$\mathbf{P}(|M_n - p| \geq 0.01) = \mathbf{P}\left(|Z_n| \geq \frac{0.01\sqrt{n}}{\sigma}\right),$$

and using the central limit theorem to approximate this last probability we write:

$$\mathbf{P}(|M_n - p| \geq 0.01) \approx \mathbf{P}\left(|Z| \geq \frac{0.01\sqrt{n}}{\sigma}\right). \quad (9)$$

where Z stands for the standard normal random variable with mean equal to 0 and variance equal to 1.

If **SOMEBODY** gives us the value of n , we would like to be able to calculate the probability $\mathbf{P}(|M_n - p| \geq 0.01)$ using the approximation in (9). But there's a slight difficulty because σ is a function that depends on p , and p is not known.

However, as we discussed in the previous section, we do know that σ is always less than or equal to $1/2$ (see figure 1 on page 3). This suggests that we could use the worst-case value of the standard deviation – replace σ by $1/2$ – and instead look at this probability:

$$\mathbf{P}(|Z| \geq 0.02\sqrt{n}). \quad (10)$$

Now, which one of the following two probabilities is larger?

$$\mathbf{P}\left(|Z| \geq \frac{0.01\sqrt{n}}{\sigma}\right) \text{ or } \mathbf{P}(|Z| \geq 0.02\sqrt{n})? \quad (11)$$

Note that since $\sigma \leq 1/2$, we have that $1/\sigma \geq 2$, and therefore

$$\frac{0.01\sqrt{n}}{\sigma} \geq 0.02\sqrt{n}.$$

Z stands for the standard normal with mean 0 and variance equal to 1:

$$Z \sim \mathcal{N}(0, 1)$$



Figure 2: The shaded area under this standard normal curve corresponds to the probability in expression (10).

It follows then that

$$\mathbf{P}\left(\left|Z\right| \geq \frac{0.01\sqrt{n}}{\sigma}\right) \leq \mathbf{P}\left(\left|Z\right| \geq 0.02\sqrt{n}\right). \quad (12)$$

This is shown in figure 3. The blue shaded area corresponds to the probability on left-hand-side of (12), and the red shaded area corresponds to the probability on the right-hand-side of the inequality.

So IF somebody gives us a value of n , we should be able to calculate the probability in (10). How do we calculate it?

The probability that the *absolute value* of Z is above $0.02\sqrt{n}$, is equal to the sum of the probabilities of the two tails that are shaded in red in figure 2. But because of the symmetry of the normal distribution, this is twice the probability of each one of the tails:

$$\mathbf{P}\left(\left|Z\right| \geq 0.02\sqrt{n}\right) = 2 \cdot (1 - \Phi(0.02\sqrt{n})).$$

Putting all of this together, we finally have a bound for the desired probability, which is expressed in terms of the standard normal CDF:

$$\mathbf{P}(|M_n - p| \geq 0.01) \leq 2 \cdot (1 - \Phi(0.02\sqrt{n})) \quad (13)$$

LET US put our finding to use and try $n = 10,000$:

$$\begin{aligned} \mathbf{P}(|M_{10,000} - p| \geq 0.01) &\leq 2 \cdot (1 - \Phi(0.02\sqrt{10,000})) \\ &= 2 \cdot (1 - \Phi(2)) \\ &= 2 \cdot (1 - 0.9772) \\ &= 0.046. \end{aligned}$$

So if we use 10,000 people in our sample, then we will get an **accuracy of one percentage point** with very high **confidence**: the probability that we do not meet the specification so that the accuracy that we get is worse than one percentage point is quite small. It's 0.046 – about 4%. This is pretty good.

SUPPOSE that your boss now tells you: "I only want the probability of not meeting the specs to be 5%."

- Specs: $\mathbf{P}(|M_n - p| \geq 0.01) \leq 0.05$

You look at your results and say: "with 10,000, I achieved a probability of a large error that's less than 5%. This means that I probably have some leeway in reducing the size of my sample."



Figure 3: The area shaded in blue is smaller than the area shaded in red under this standard normal curve. The blue area is the probability on the left-hand-side of (12). The red area is the probability on the right-hand-side of (12).

0.8	.7881	.7910	.7939	.7967	.7995	.8023
0.9	.8159	.8186	.8212	.8238	.8264	.8289
1.0	.8413	.8438	.8461	.8485	.8508	.8531
1.1	.8643	.8665	.8686	.8708	.8729	.8749
1.2	.8849	.8869	.8888	.8907	.8925	.8944
1.3	.9032	.9049	.9066	.9082	.9099	.9115
1.4	.9192	.9207	.9222	.9236	.9251	.9265
1.5	.9332	.9345	.9357	.9370	.9382	.9394
1.6	.9452	.9463	.9474	.9484	.9495	.9505
1.7	.9554	.9564	.9573	.9582	.9591	.9599
1.8	.9641	.9649	.9656	.9664	.9671	.9678
1.9	.9713	.9719	.9726	.9732	.9738	.9744
2.0	.9772	.9778	.9783	.9788	.9793	.9798

Figure 4: A fragment of the standard normal table. We are looking for the value of the CDF at point 2.

What could the size of the sample be and still meet those specs?

What we're trying to do here is that we have (13), the approximation for the probability of interest, and we want to set this probability to a value of 0.05:

$$P(|M_n - p| \geq 0.01) \leq 2 \cdot (1 - \Phi(0.02\sqrt{n})) = 0.05 \quad (14)$$

Then we want to ask, what is the value of n that will result in the probability of not meeting the specs, which is 0.05?

So we do the algebra and we find that (14) corresponds to requiring that $\Phi(0.02\sqrt{n}) = 0.975$. What's the interpretation of this? We want to choose n so that the probability of the two green tails in figure 5 is 5%. This means that we want the probability of one of the tails shown in figure 4 to be 2.5%.

So we peek into the standard normal table (or rush to our favorite software package) in hopes of finding the value for which the CDF is equal to 0.975. We look around and we find that 0.975 corresponds to 1.96.

This tells us that $0.02\sqrt{n} = 1.96$, which we solve for n to find that $n = 9,604$. This is indeed some reduction from the 10,000 that we had originally.

How DOES THIS relate to the real world? When you read newspapers or check websites about polls, you will never see sample sizes that are about 10,000. You will usually see sample sizes of the order of 1,000, sometimes even smaller.

How can they do that? Well, they can do that because the specs that they impose are not as tight as the specs that we have here. Usually, they tell you that the results are accurate within three percentage points, let's say, instead of one percentage point. And by moving from 0.01 to 0.03, if you repeat the calculations, you will find that the sample size of about 1,000 will actually do.



Figure 5: The total area shaded in green corresponds to the probability of 0.05. Therefore, because of symmetry, each tail has the probability of 0.025. This implies that $\Phi(0.02\sqrt{n})$ must be 0.975.

0.8	.7881	.7910	.7939	.7967	.7995	.8023	.8051	↓
0.9	.8159	.8186	.8212	.8238	.8264	.8289	.8315	↓
1.0	.8413	.8438	.8461	.8485	.8508	.8531	.8554	↓
1.1	.8643	.8665	.8686	.8708	.8729	.8749	.8770	↓
1.2	.8849	.8869	.8888	.8907	.8925	.8944	.8962	↓
1.3	.9032	.9049	.9066	.9082	.9099	.9115	.9131	↓
1.4	.9192	.9207	.9222	.9236	.9251	.9265	.9279	↓
1.5	.9332	.9345	.9357	.9370	.9382	.9394	.9406	↓
1.6	.9452	.9463	.9474	.9484	.9495	.9505	.9515	↓
1.7	.9554	.9564	.9573	.9582	.9591	.9599	.9608	↓
1.8	.9641	.9649	.9656	.9664	.9671	.9678	.9686	↓
1.9	.9713	.9719	.9726	.9732	.9738	.9744	.9750	↓
2.0	.9772	.9778	.9783	.9788	.9793	.9798	.9803	↓

Figure 6: A fragment of the standard normal table. We are looking for the value of the standard normal at which its CDF is equal to 0.975. This happens at point 1.96.